

WIDS ENDTERM REPORT

I am Yash Goutharia(25B1004). A first year student who was allotted the project Intro to reinforcement learning. Reinforcement learning is a learning framework in which an agent interacts with an environment by taking actions, receives scalar rewards, and learns a policy that maximizes expected cumulative reward over time through trial and error, without being told the correct actions. For me, involvement in this project gave me a foundation to learn how a system can interact with its environment, similar to how humans learn from experience. This idea felt very natural and exciting to me.

For this WIDS, we mainly studied the books Sutton & Barto: Reinforcement Learning & Grokking: Deep Reinforcement Learning. Also some video tutorials and content for learning python was provided by the mentor.

Week 1

This week was focused on understanding value function and policies, and implementing algorithms on the Multi-Armed Bandit Problem


Policy:- A policy π is a mapping from states to a probability distribution over actions:

Value function:- The value function of a policy π is the expected return starting from state s and thereafter following π


The multi-armed bandit problem models an agent facing a set of k actions (arms), each producing stochastic rewards from a Normal distribution with unknown mean. At each time step, the agent chooses one arm and observes its reward, with the objective of maximizing cumulative reward over time. In our version, there are 10 arms.

We implemented the following algorithms


- 1.) Greedy Algorithm:- Always selects the action with the highest estimated value so far. It exploits current knowledge but does **no exploration**




2.) Epsilon Greedy Algorithm:- With probability $1-\epsilon$, chooses the greedy (best-known) action, and with probability ϵ , selects a random action. This has both exploration and exploitation



3.) Optimistic Initial Values:- Initialize all action-value estimates to unrealistically high values. Early actions tend to look bad after being tried once, which naturally encourages exploration without explicit randomness



4.) Upper Confidence Bound:- Selects actions based on both estimated value and uncertainty, choosing the action that maximizes $Q(a) + c \cdot \sqrt{\ln(t) / N(a)}$



Week 2

This week consisted of understanding Markov decision processes and coding them and also understanding Bellman Optimality Equations

A Markov Decision Process (MDP) is a mathematical framework used to model sequential decision-making where outcomes are partly random and partly controlled by an agent. MDPs define an agent interacting with an environment over discrete time steps to maximize long-term rewards, adhering to the Markov property where future states depend only on the current state and action. An MDP consists of the 5 parts


- 1.) S = Set of all possible states
- 2.) A = Set of all possible actions
- 3.) P = transition probability function
- 4.) R = reward function
- 5.) γ = discount factor

After this, we coded in various sample MDPs, which were

1.) Slippery Walk

The Slippery Walk Five environment

(1) This environment is stochastic, and even if the agent selects the Right action, there's a chance it goes left!




The agent starts in *S*, *H* is a hole, *G* is the goal and provides a +1 reward.

2.) Frozen Lake

The frozen lake (FL) environment

(1) Agent starts each trial here.

(4) But, these are holes that, if the agent falls into, will end the episode right away.




(2) Note that the slippery, frozen surface may send the agent to unintended places.

(3) Agent gets a +1 when it arrives here.

3.) A deterministic version of the frozen Lake, but with a lock on the goal state and a key, and so slightly different state functions.

We also made an MDP for a game called Lights Out



Lights Out is an electronic game released by Tiger Electronics in 1995. The game consists of a 5 by 5 grid of lights. When the game starts, a random number or a stored pattern of these lights is switched on. Pressing any of the lights will toggle it and the adjacent lights. The goal of the puzzle is to switch all the lights off, preferably with as few button presses as possible.

We created a 4x4 version of Lights Out with the following attributes

- The states should represent the game grid layout
- The actions are the cells you wish to toggle
- Give a reward of 1 when solved
- The ONLY terminal state is the fully solved state

We also studied Bellman optimality equations, which allow us to iteratively compute the optimal value functions.

Bellman Equation for value functions,

Bellman Equation for v_π :

$$v_\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(a|s) \sum_{\substack{s' \in \mathcal{S} \\ r \in \mathcal{R}}} p(s', r|s, a) [r + \gamma v_\pi(s')]$$

Bellman Equation for action value functions,

Bellman Equation for q_π :

$$q_\pi(s, a) = \sum_{\substack{s' \in \mathcal{S} \\ r \in \mathcal{R}}} p(s', r | s, a) \left[r + \gamma \sum_{a' \in \mathcal{A}(s')} \pi(a' | s') q_\pi(s', a') \right]$$

Week 3

This week consisted of learning about policy iteration and Value iteration and solving two problems from Sutton & Barto and solving the Lights out game made previously, The first problem we had to solve was Jack's Car rentals which is formulated as

1.) Jack manages two locations for a nationwide car rental company. Each day, some number of customers arrive at each location to rent cars. If Jack has a car available, he rents it out and is credited \$10 by the national company. If he is out of cars at that location, then the business is lost. Cars become available for renting the day after they are returned. To help ensure that cars are available where they are needed, Jack can move them between the two locations overnight, at a cost of \$2 per car moved. We assume that the number of cars requested and returned at each location are Poisson random variables, meaning that the probability that the number is n is $(\lambda)^n \cdot e^{-(\lambda)} / n!$, where λ is the expected number. Suppose λ


3 and 4 for rental requests at the first and second locations and 3 and 2 for returns. To simplify the problem slightly, we assume that there can be no more than 20 cars at each location (any additional cars are returned to the nationwide company, and thus disappear from the problem) and a maximum of five cars can be moved from one location to the other in one night.

Now, this we solved using policy iteration, which is a combination of policy evaluation, where you use the policy to make a better value function, and policy improvement, where you use the value function to generate a better policy

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization
 $V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$
2. Policy Evaluation
 Loop:
 $\Delta \leftarrow 0$
 Loop for each $s \in \mathcal{S}$:
 $v \leftarrow V(s)$
 $V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$
 $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
 until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)
3. Policy Improvement
 $policy-stable \leftarrow true$
 For each $s \in \mathcal{S}$:
 $old-action \leftarrow \pi(s)$
 $\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$
 If $old-action \neq \pi(s)$, then $policy-stable \leftarrow false$
 If $policy-stable$, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

My solution to this looked like



The second problem we had to solve was the Gambler's problem which is formulated as

2.) A gambler has the opportunity to make bets on the outcomes of a sequence of coin flips. If the coin comes up heads, he wins as many dollars as he has staked on that flip; if it is tails, he loses his stake. The game ends when the gambler wins by reaching his goal of \$100, or loses by running out of money. On each flip, the gambler must decide what portion of his capital to stake, in integer numbers of dollars.

For this, we had to implement Value iteration, which involves repeatedly generating more accurate value functions. Until Value function converges and then using that value function to create a policy.

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation
Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

| $\Delta \leftarrow 0$

| Loop for each $s \in \mathcal{S}$:

| $v \leftarrow V(s)$

| $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$


| $\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

Output a deterministic policy, $\pi \approx \pi_*$, such that

$\pi(s) = \arg\max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$

So the solution I made looks like




Then we used the MDP from the previous week and used it to try to find the optimal way to solve Lights Out

The graph is not unique but follows almost the same pattern of capital vs stake

Week 4

For the final week, we made a RL based solver for the 15 puzzle.

The **15 puzzle** is a sliding puzzle. It has 15 square tiles numbered 1 to 15 in a frame that is 4 tile positions high and 4 tile positions wide, with one unoccupied position. Tiles in the same row or column of the open position can be moved by sliding them horizontally or vertically, respectively. The goal of the puzzle is to place the tiles in numerical order (from left to right, top to bottom).



To solve it we drew inspiration from the ideas presented in the article

<https://medium.com/@amshali/15-puzzle-with-reinforcement-learning-8bcfc1aa54e7>

I have solved the problem in python and tried to implement it using a hierarchical rl algorithm as mentioned in the article. It inculcates solving the problem row wise starting from the last row.

```
GOAL = np.array([
    [1, 2, 3, 4],
    [5, 6, 7, 8],
    [9, 10, 11, 12],
    [13, 14, 15, 0]
])
```